

# Machine Learning

## Lecture 2: Linear Regression

Feng Li

[fli@sdu.edu.cn](mailto:fli@sdu.edu.cn)

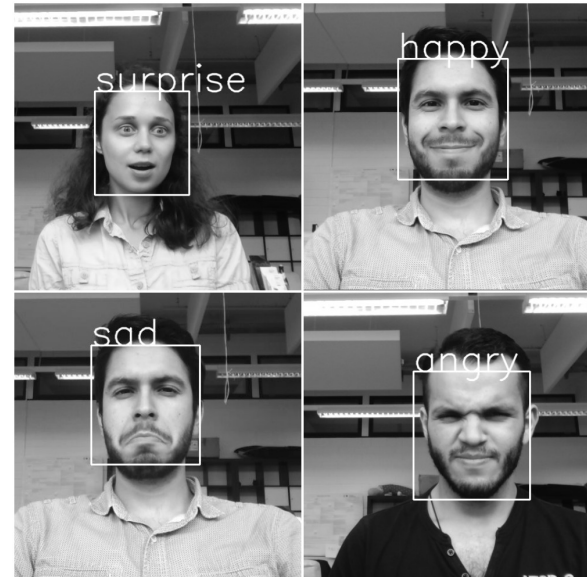
<https://funglee.github.io>

School of Computer Science and Technology  
Shandong University

# Supervised Learning

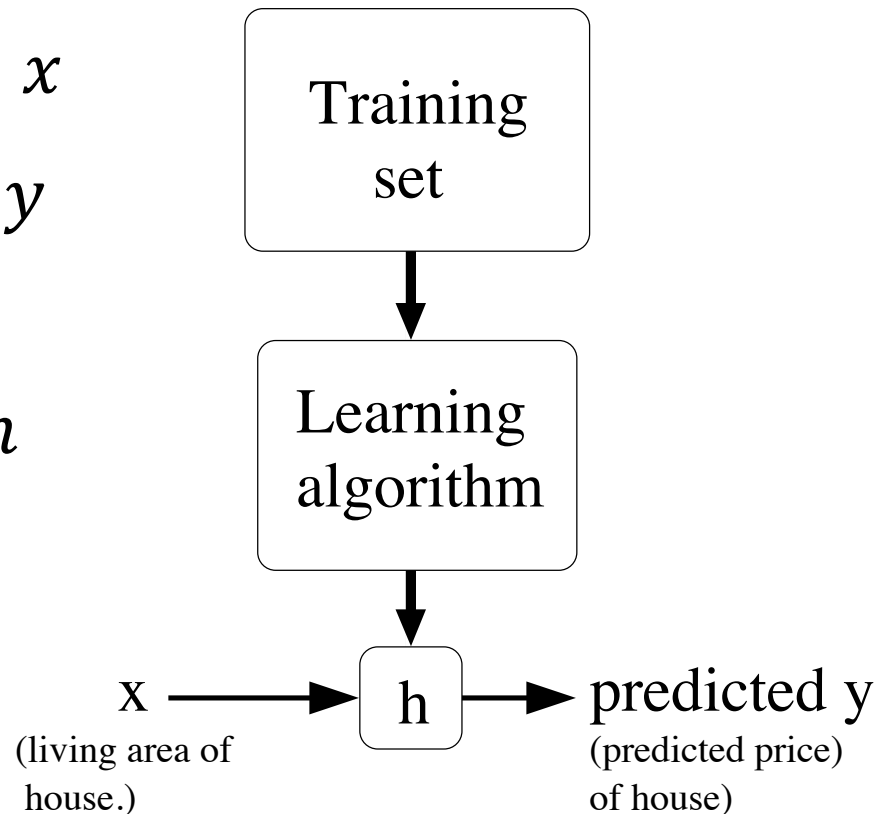
- Regression: Predict a continuous value
- Classification: Predict a discrete value, the class

Living area (feet <sup>2</sup> )	Price (1000\$s)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮



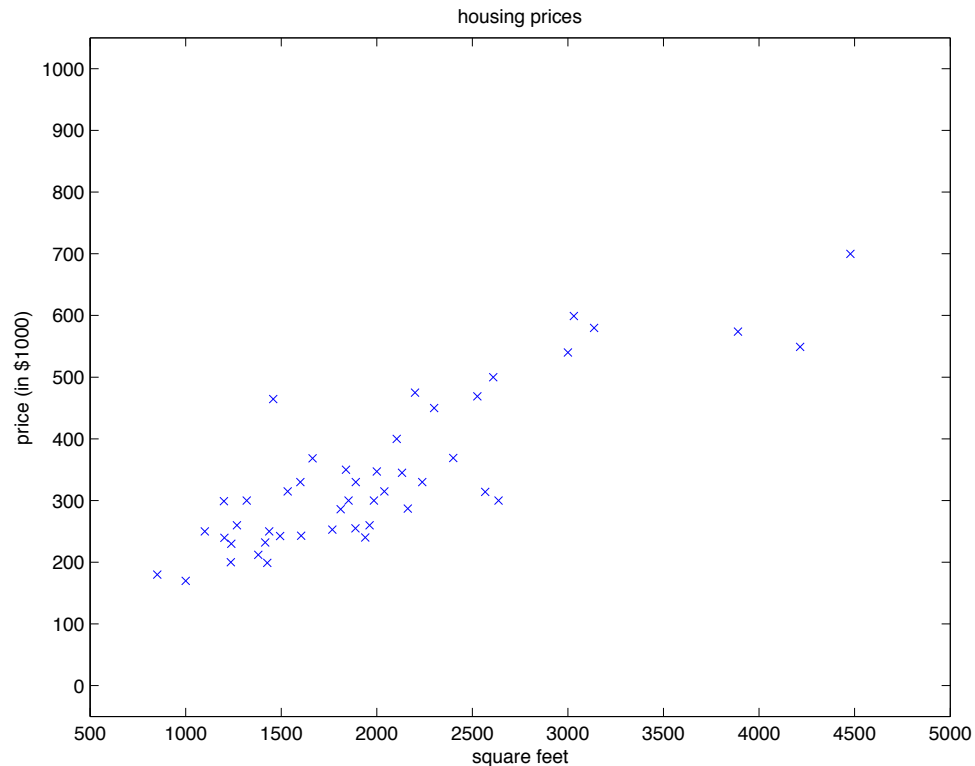
# Supervised Learning (Contd.)

- Features: Input variables,  $x$
- Target: Output variables,  $y$
- Training examples:  
 $(x^{(i)}, y^{(i)}), i = 1, 2, \dots, m$
- Hypothesis:  $h: \mathcal{X} \rightarrow \mathcal{Y}$



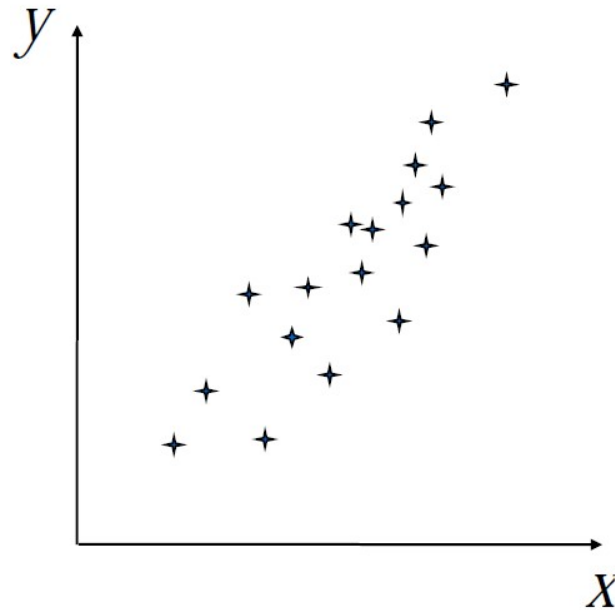
# Linear Regression

- Linear hypothesis:  $h(x) = \theta_1 x + \theta_0$
- $\theta_i$  ( $i = 1, 2$  for 2D cases): Parameters to estimate
- How to choose  $\theta_i$  's ?



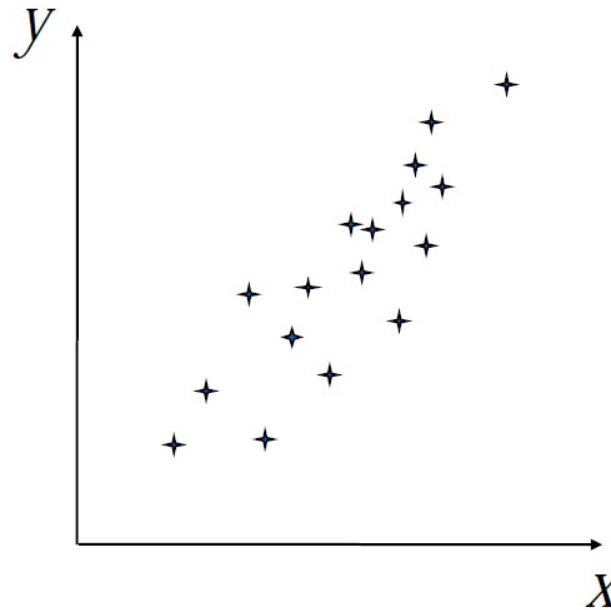
# Linear Regression (Contd.)

- Input: Training set  $(x^{(i)}, y^{(i)}) \in \mathbb{R}^2$  with  $i = 1, 2, \dots, m$
- Goal: Model the relationship between  $x$  and  $y$  such that we can predict the target  $y = h(x)$  according to a given new feature input  $x$



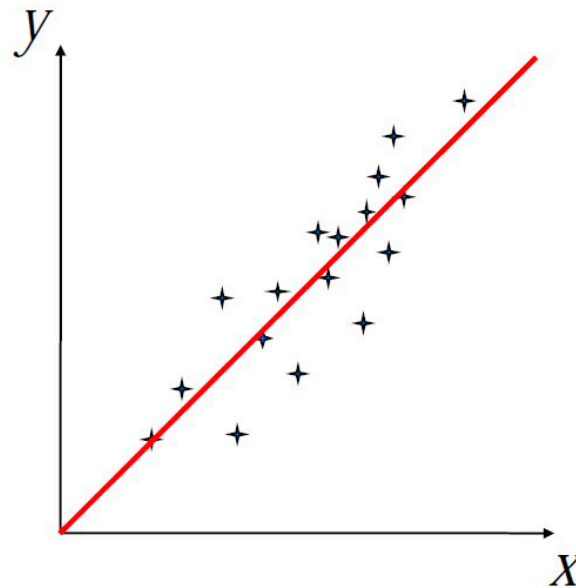
# Linear Regression (Contd.)

- Input: Training set  $(x^{(i)}, y^{(i)}) \in \mathbb{R}^2$  with  $i = 1, 2, \dots, m$
- Goal: Model the relationship between  $x$  and  $y$  such that we can predict the target  $y = h(x)$  according to a given new feature input  $x$



# Linear Regression (Contd.)

- The relationship between  $x$  and  $y$  is modeled as a linear function (with respect to  $\theta$ ).
- The linear function in the 2D plane is a straight line
- Hypothesis:  $h_{\theta}(x) = \theta_0 + \theta_1 x$



# Linear Regression (Contd.)

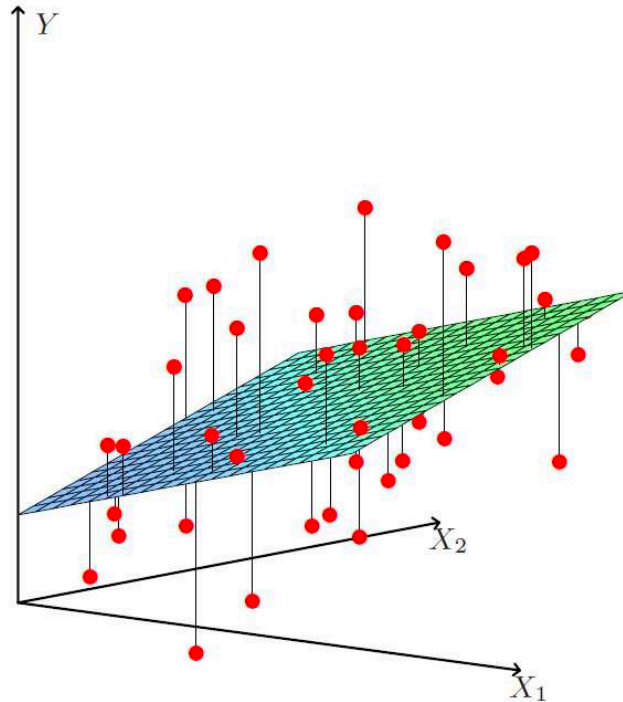
- Given data  $x \in \mathbb{R}^n$ , we then have  $\theta \in \mathbb{R}^{n+1}$
- Hence,  $h_{\theta}(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$ , where  $x_0 = 1$
- What is the best choice of  $\theta$

$$\min_{\theta} J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

where  $J(\theta)$  is the so-called *cost function*



# Linear Regression (Contd.)



$$\min_{\theta} J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

# Gradient Descent (GD) Algorithm

- If the multi-variable function  $J(\theta)$  is differentiable in a neighborhood of a point  $\theta$ , then  $J(\theta)$  decreases fastest if one goes from  $\theta$  in the direction of the negative gradient of  $J$  at  $\theta$
- Find a local minimum of a differentiable function using gradient descent

---

**Algorithm 1** Gradient Descent

---

- 1: **Given** a starting point  $\theta \in \text{dom } J$
  - 2: **Repeat**
  - 3:   Calculate gradient  $\nabla J(\theta)$
  - 4:   Update  $\theta \leftarrow \theta - \alpha \nabla J(\theta)$
  - 5: **until** convergence criterion is satisfied
- 

Remarks:  $\theta$  is usually initialized randomly, and  $\alpha$  is so-called *learning rate*

# GD Algorithm (Contd.)

- Stopping criterion
  - The gradient has its magnitude less than or equal to a predefined threshold (say  $\varepsilon$ ), i.e.,

$$\|\nabla f(x)\|_2 \leq \varepsilon$$

where  $\|\cdot\|_2$  is  $\ell_2$  norm, such that the values of the objective function differ very slightly across different iterations

- Set a fixed value for the maximum number of iterations, such that the algorithm is terminated after the number of the iterations exceeds the threshold

# GD Algorithm (Contd.)

- Specifically, we update each component of  $\theta$  according to the following rule

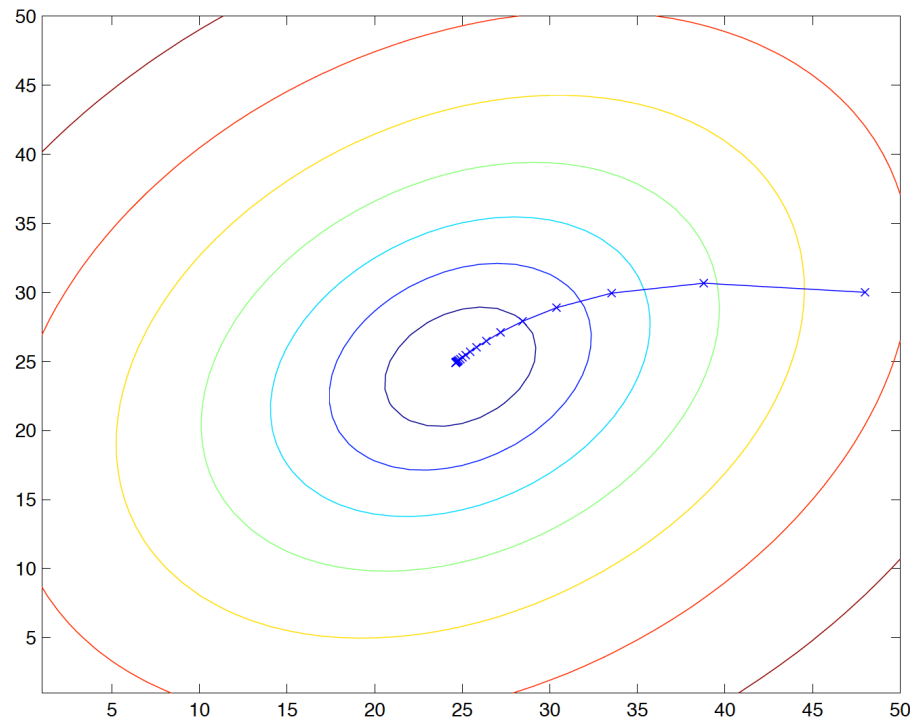
$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}, \quad \forall j$$

- Calculating the gradient for linear regression

$$\begin{aligned} \frac{\partial J(\theta)}{\partial \theta_j} &= \frac{\partial}{\partial \theta_j} \frac{1}{2} \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)})^2 \\ &= \sum_{i=1}^m (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)} \end{aligned}$$

# GD Algorithm (Contd.)

- An illustration of gradient descent algorithm



The objective function is decreased along the gradient

# GD Algorithm (Contd.)

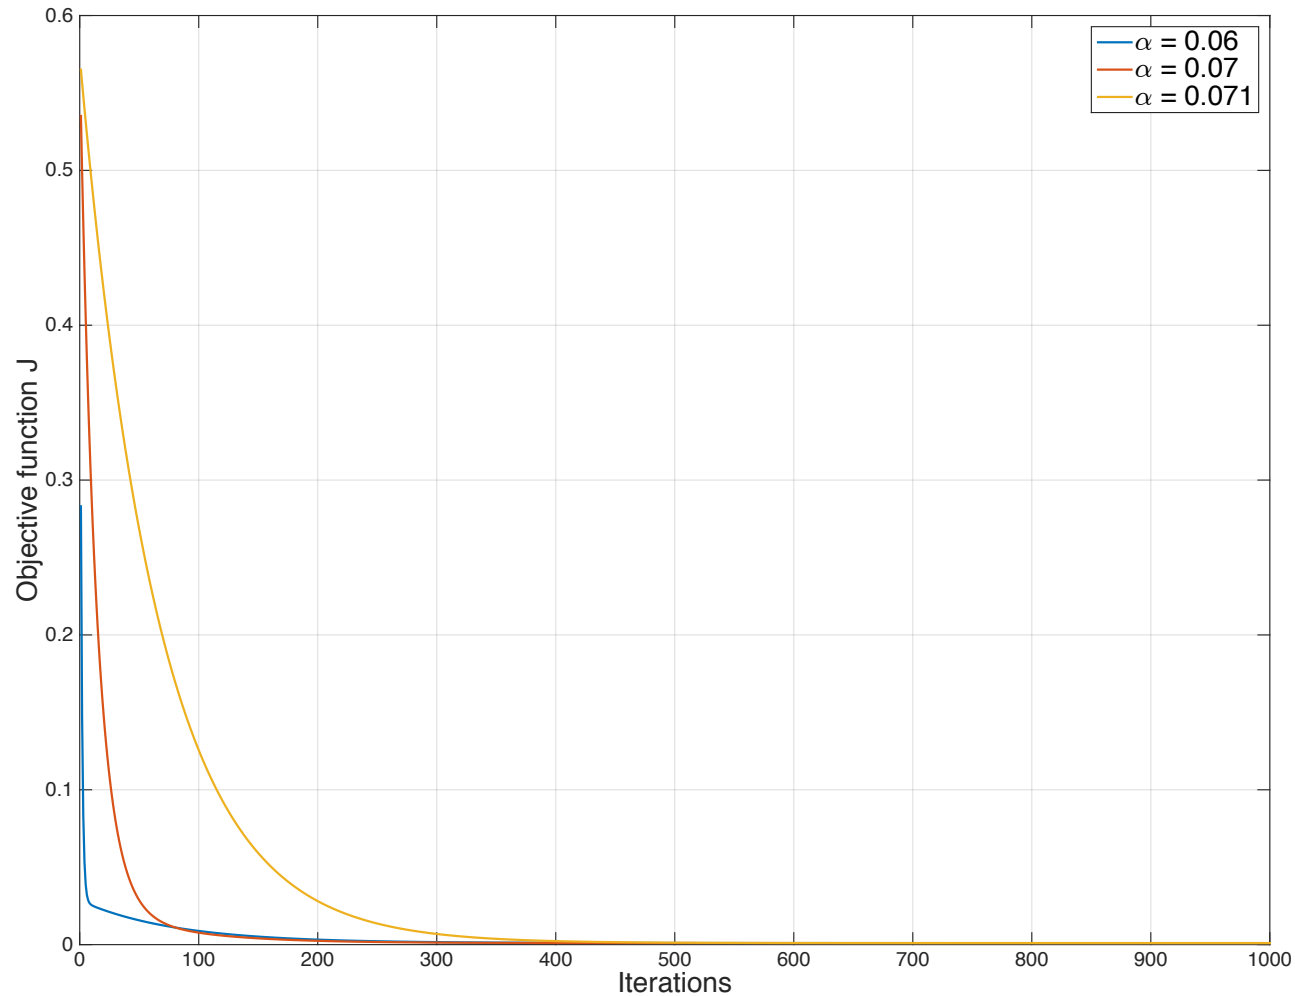
- Another commonly used form

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Gradient ascent algorithm
  - Maximize the differentiable function  $J(\theta)$
  - The gradient represents the direction along which  $J$  increase fastest
  - Therefore, we have

$$\theta_j \leftarrow \theta_j + \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$

# Convergence under Different Step Sizes



# GD Algorithm (Contd.)

## ➤ What if the training set is huge?

- In the above batch gradient descent algorithm, we have to run through the entire training set in each iteration
- A considerable computation cost is induced!

## ➤ Stochastic gradient descent (SGD), also known as incremental gradient descent, is a stochastic approximation of the gradient descent optimization method

- In each iteration, the parameters are updated according to the gradient of the error with respect to one training sample only

---

**Algorithm 2** Stochastic Gradient Descent for Linear Regression

---

```
1: Given a starting point  $\theta \in \text{dom } J$ 
2: repeat
3:   Randomly shuffle the training data;
4:   for  $i = 1, 2, \dots, m$  do
5:      $\theta \leftarrow \theta - \alpha \nabla J(\theta; x^{(i)}, y^{(i)})$ 
6:   end for
7: until convergence criterion is satisfied
```

---



# Stochastic Gradient Descent (SGD)

➤ Stochastic gradient descent (SGD), also known as incremental gradient descent, is a stochastic approximation of the gradient descent optimization method

- In each iteration, the parameters are updated according to the gradient of the error with respect to one training sample only
- For linear regression,

$$\nabla J(\theta; x^{(i)}, y^{(i)}) = (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)}$$

---

**Algorithm 2** Stochastic Gradient Descent for Linear Regression

---

- 1: **Given** a starting point  $\theta \in \text{dom } J$
  - 2: **repeat**
  - 3:   Randomly shuffle the training data;
  - 4:   **for**  $i = 1, 2, \dots, m$  **do**
  - 5:      $\theta \leftarrow \theta - \alpha \nabla J(\theta; x^{(i)}, y^{(i)})$
  - 6:   **end for**
  - 7: **until** convergence criterion is satisfied
-

# SGD (Contd.)

- Stochastic gradient descent (SGD), also known as incremental gradient descent, is a stochastic approximation of the gradient descent optimization method
- In each iteration, the parameters are updated according to the gradient of the error with respect to one training sample only
  - For linear regression,

$$\nabla J(\theta; x^{(i)}, y^{(i)}) = (\theta^T x^{(i)} - y^{(i)}) x_j^{(i)}$$

---

**Algorithm 2** Stochastic Gradient Descent for Linear Regression

---

- 1: **Given** a starting point  $\theta \in \text{dom } J$
  - 2: **repeat**
  - 3:   Randomly shuffle the training data;
  - 4:   **for**  $i = 1, 2, \dots, m$  **do**
  - 5:      $\theta \leftarrow \theta - \alpha \nabla J(\theta; x^{(i)}, y^{(i)})$
  - 6:   **end for**
  - 7: **until** convergence criterion is satisfied
-

# More about SGD

- The objective does not always decrease for each iteration
- Usually, SGD has approaching the minimum much faster than batch GD
- SGD may never converge to the minimum, and oscillating may happen
- A variants: Mini-batch, say pick up a small group of samples and do average, which may accelerate and smoothen the convergence

# Matrix Derivatives<sup>1</sup>

- A function  $f: \mathbb{R}^{m \times n} \rightarrow \mathbb{R}$
- The derivative of  $f$  with respect to  $A$  is defined as

$$\nabla f(A) = \begin{bmatrix} \frac{\partial f}{\partial A_{11}} & \cdots & \frac{\partial f}{\partial A_{1n}} \\ \vdots & \ddots & \vdots \\ \frac{\partial f}{\partial A_{m1}} & \cdots & \frac{\partial f}{\partial A_{mn}} \end{bmatrix}$$

- For an  $n \times n$  matrix, its trace is defined as  $\text{tr}A = \sum_{i=1}^n A_{ii}$

$$\text{tr}ABCD = \text{tr}DABC = \text{tr}CDAB = \text{tr}BCDA$$

$$\text{tr}A = \text{tr}A^T, \text{tr}(A + B) = \text{tr}A + \text{tr}B, \text{tr}aA = a\text{tr}A$$

$$\nabla_A \text{tr}AB = B^T, \nabla_{A^T} f(A) = (\nabla_A f(A))^T$$

$$\nabla_A \text{tr}ABA^T C = CAB + C^T AB^T, \nabla_A |A| = |A|(A^{-1})^T$$

$$\text{Funky trace derivative } \nabla_{A^T} \text{tr}ABA^T C = B^T A^T C^T + BA^T C$$

# Revisiting Least Square

- Assume

$$X = \begin{bmatrix} (x^{(1)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix}, \quad Y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

- Therefore, we have

$$X\theta - Y = \begin{bmatrix} (x^{(1)})^T \theta \\ \vdots \\ (x^{(m)})^T \theta \end{bmatrix} - \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix} = \begin{bmatrix} h_{\theta}(x^{(1)}) - y^{(1)} \\ \vdots \\ h_{\theta}(x^{(m)}) - y^{(m)} \end{bmatrix}$$

- The objective function can be written as

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2} (X\theta - Y)^T (X\theta - Y)$$

# Revisiting Least Square (Contd.)

- Minimize  $J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2} (X\theta - Y)^T (X\theta - Y)$
- Calculate its derivative with respect to  $\theta$

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \frac{1}{2} (Y - X\theta)^T (Y - X\theta) \\&= \frac{1}{2} \nabla_{\theta} (Y^T - \theta^T X^T) (Y - X\theta) \\&= \frac{1}{2} \nabla_{\theta} \text{tr}(Y^T Y - Y^T X\theta - \theta^T X^T Y + \theta^T X^T X\theta) \\&= \frac{1}{2} \nabla_{\theta} \text{tr}(\theta^T X^T X\theta) - X^T Y \\&= \frac{1}{2} (X^T X\theta + X^T X\theta) - X^T Y \\&= X^T X\theta - X^T Y\end{aligned}$$

# Revisiting Least Square (Contd.)

- **Theorem:**

The matrix  $A^T A$  is invertible if and only if the columns of  $A$  are linearly independent. In this case, there exists only one least-squares solution

$$\theta = (X^T X)^{-1} X^T Y$$

- Prove the above theorem in Problem Set 1

# Probabilistic Interpretation

- The target variables and the inputs are related

$$y = \theta^T x + \epsilon$$

- $\epsilon$ 's denote the errors and are independently and identically distributed (i.i.d.) according to a Gaussian distribution  $\mathcal{N}(0, \sigma^2)$

- The density of  $\epsilon^{(i)}$  is given by

$$f(\epsilon) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\epsilon^2}{2\sigma^2}\right)$$

- The conditional probability density function of  $y$

$$y \mid x; \theta \sim \mathcal{N}(\theta^T x, \sigma^2)$$



# Probabilistic Interpretation (Contd.)

- The training data  $\{x^{(i)}, y^{(i)}\}_{i=1, \dots, m}$  are sampled identically and independently

$$p(y^{(i)} | x^{(i)}; \theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right)$$

- Likelihood function

$$L(\theta) = \prod_i p(y^{(i)} | x^{(i)}; \theta)$$

$$= \prod_i \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right)$$

# Probabilistic Interpretation (Contd.)

- Maximizing the likelihood  $L(\theta)$ 
  - Choosing the optimal  $\theta$  to make the data as high probability as possible
- Since  $L(\theta)$  is complicated, we maximize an logarithmic function of  $L(\theta)$  instead

$$\begin{aligned}\ell(\theta) &= \log L(\theta) \\ &= \log \prod_i^m \frac{1}{\sqrt{2\pi}\sigma} \exp \left( -\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2} \right) \\ &= \sum_i^m \log \frac{1}{\sqrt{2\pi}\sigma} \exp \left( -\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2} \right) \\ &= m \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{2\sigma^2} \sum_i (y^{(i)} - \theta^T x^{(i)})^2\end{aligned}$$

Thanks !

Q & A